Project Charter

AI Customer Support Platform - authorisation, objectives, scope summary and governance

Document	01 - Project Charter
Project	AI Customer Support Platform (fictional company: Kestrel Commerce Ltd.)
Prepared by	Project Manager (portfolio author) - self-directed simulation
Status	Baseline v1.0 - simulated approval 7 Jan 2026 (M1)
Project window	5 Jan - 13 Mar 2026 (simulated project window)

SIMULATED PROJECT - Self-directed portfolio case study. Not professional experience, not client work, not employment. All people, figures and outcomes are fictional.

1. Project overview

Project name	AI Customer Support Platform (codename KestrelAssist)
Organisation	Kestrel Commerce Ltd. (fictional)
Executive Sponsor	Marcus Hale, VP Customer Experience (fictional)
Product Owner	Elena Rossi, Product Manager (fictional)
Project Manager	Project Manager (portfolio author)
Methodology	Scrum, 5 two-week sprints (proposed delivery methodology in a self-directed simulation)
Planned dates	5 Jan 2026 to 13 Mar 2026 (10 weeks); pilot release 11 Mar 2026
Team	8 roles, 61 person-weeks of planned capacity (Project Manager, Product Manager, AI/ML, Backend, Frontend, UX, QA, DevOps)

2. Business case

Kestrel's support team handles about 9,000 tickets a month with a 6.5-hour average first response (baseline assumptions). Roughly 45% of tickets are routine questions that the knowledge base already answers. An AI assistant that answers from approved content, shows its sources and escalates when unsure can remove much of that routine load, speed up responses and give managers data on content gaps.

Options considered: (1) do nothing, (2) hire more agents, (3) improve help-center search only, (4) AI assistant grounded in the knowledge base with human escalation. Option 4 was selected for the MVP because it addresses speed and workload together, and its main risk (wrong answers) can be controlled with grounding, fallback and evaluation gates.

Estimated cost (illustrative planning assumptions)

Item	Amount (USD)
Cloud hosting (3 environments, managed DB, monitoring)	3,000
LLM + embedding API usage (build, evaluation and 4-week pilot)	1,500
Transactional email service and tooling licences	500
Contingency (10% of non-labour budget)	500
Total non-labour budget	5,500

Labour is funded from existing internal team capacity (61 person-weeks planned) and is not separately costed in this simulation.

Expected benefits (Proposed Project Targets, measured after go-live)

- At least 35% of conversations resolved without handover within 8 weeks of pilot go-live (K1).
- AI answers in 8 seconds or less (p95); first human reply to escalations within 2 business hours (K2).
- At least 75% positive feedback on AI answers (K3); at least 80% of the 20 most frequent ticket categories covered by a reviewed article (K4).

3. Problem statement, goal and objectives

Problem statement	Merchants with routine questions wait about 6.5 hours (baseline assumption) for answers that already exist in the knowledge base, while agents spend a large share of their time on repetitive tickets.
Project goal	Deliver a pilot-ready AI customer support platform (MVP) in 10 weeks, within a 5,500 USD non-labour budget, that answers from the knowledge base, escalates to humans, and gives managers basic analytics.
Objectives	O1 Grounded AI answers with citations. O2 Safe handover to human agents with full context. O3 Admin-managed knowledge base. O4 Basic analytics for managers. O5 Authentication, security and privacy controls suitable for a pilot. O6 Release by 11 Mar 2026.

4. Scope summary

In scope (MVP)	Out of scope (this release)
AI chat from knowledge base with citations and fallback Human escalation, agent queue and reply Knowledge-base admin (create, edit, publish, bulk import) Conversation history (30 days) Basic analytics dashboard Staff sign-in, customer email-code sign-in, 4 roles Security and privacy controls (PII redaction, retention, deletion) UAT, performance test, pilot release, guides and training	Helpdesk / CRM integration Multiple languages Native mobile app, voice or social channels AI performing account actions (refunds, plan changes) Single sign-on with the merchant portal Advanced analytics / BI exports Fine-tuning or training custom models Full-scale rollout beyond the pilot cohort

Full detail: 02_Project_Scope.pdf. Requirements: 03_Requirements.pdf.

5. Major deliverables

Deliverable	Description	Milestone
Approved charter and requirements baseline	Charter, scope, requirements, backlog and UX prototype	M1, M2
Knowledge base and AI core	Admin KB tools, ingestion pipeline, RAG answer API, guardrails, golden set	M3
Customer chat experience	Sign-in, chat with sources, feedback, history, escalation request, PII redaction	M4
Human handover and analytics	Agent queue and console, customer reply, event tracking, manager dashboard, hardening	M5
Validated release	UAT sign-off, load test report, runbook, guides, training	M6, M7
Pilot platform in production	Production release with monitoring, rollback and hypercare	M8
Closure pack	Closure report, lessons learned, handover, post-MVP backlog	M9

6. Key stakeholders

Stakeholder	Role	Strategy
Marcus Hale - VP Customer Experience	Executive Sponsor; funds the project and decides Go/No-Go	Manage closely
Elena Rossi - Product Manager (Product Owner)	Owns backlog, priorities and acceptance of stories	Manage closely
Sofia Lindqvist - Support Manager	Business owner of support operations; UAT lead and approver	Manage closely
Raj Patel - CTO / Head of Engineering	Approves architecture, security stance and engineering capacity	Keep satisfied
Hannah Weber - Security & Privacy Officer	Reviews privacy and security controls; go-live gatekeeper	Keep satisfied
Daniel Moreau - Customer Success Lead	Consumer of analytics; represents merchant voice; raised CR-001	Keep informed
Aiden Park - Knowledge & Content Lead (KB Admin)	Admin end user; owns KB content quality	Keep informed
Jamal Brooks and the 14 support agents	Agent end users; usability testers and UAT participants	Keep informed

Full register with engagement approach: 09_Stakeholder_Register.xlsx.

7. Assumptions and constraints

Assumptions

- The team (8 roles) is allocated as planned from 5 Jan 2026; part-time roles keep their allocation.
- A hosted LLM provider is available on acceptable data-handling terms (to be confirmed by Legal; risk R04).
- Support SMEs can give about 2-3 hours per week to content cleanup and reference answers.
- About 320 legacy help-center articles can be exported; roughly 280 remain after cleanup.
- English only; pilot cohort of merchants selected by Customer Success.
- Cloud account and budget are released in Week 1; no public holidays are assumed in the 10 weeks.

Constraints

- Fixed duration: 10 weeks, pilot release on 11 Mar 2026.
- Non-labour budget capped at 5,500 USD including 10% contingency.
- Team capacity fixed at 61 person-weeks; no additional hires.
- Customer data may only be sent to an LLM provider after redaction and confirmed provider terms.
- The AI may answer only from published knowledge-base content.
- MVP must run as a standalone product with no integration to existing helpdesk systems.

8. Major risks (top six by score)

ID	Risk	Score	Main mitigation
R01	AI hallucination: the assistant states something not supported by the knowledge base (for example a wrong billing rule).	20	Answer only from retrieved passages with citations (AIR-01).
R03	Poor knowledge-base content: duplicate, outdated or conflicting articles produce wrong answers.	16	Content audit and cleanup before indexing (WBS 6.3) with Support SMEs.
R08	Scope creep: new feature requests (languages, integrations, extra channels) after the baseline.	16	Baselined scope with explicit Won't-Have list.
R04	Data privacy: personal data is sent to the LLM provider, retained too long, or cannot be deleted on request.	15	PII redaction before every LLM call (US-25).
R05	Security vulnerabilities: broken access control, injection or prompt injection exposing data or tampering with answers.	15	RBAC (US-05).
R02	Incorrect or irrelevant answers because retrieval picks the wrong passages (ambiguous questions, poor chunking).	12	Chunking and ranking tuned in Sprint 2-3.

Full register with 17 risks: 11_Risk_Register.xlsx.

9. Milestones

ID	Milestone	Date
M1	Project Charter approved	7 Jan 2026
M2	Requirements & UX baseline approved	16 Jan 2026
M3	Sprint 2 review: KB ingestion + RAG answers demo	30 Jan 2026
M4	Sprint 3 review: customer chat + escalation request demo	13 Feb 2026
M5	Feature complete / code freeze	27 Feb 2026
M6	UAT sign-off	6 Mar 2026
M7	Go / No-Go decision	9 Mar 2026
M8	Pilot release to production	11 Mar 2026
M9	Project closure	13 Mar 2026

10. Proposed success criteria

All numbers are Proposed Project Targets for this simulation, not actual results.

ID	Go-live criterion	Proposed Project Target
G1	Golden-set answer accuracy	≥ 85%
G2	Unsupported (hallucinated) answers on golden + adversarial sets	≤ 5%
G3	Unanswerable questions that trigger fallback	≥ 95%
G4	p95 response time at 50 concurrent sessions	≤ 8 s, error rate < 1%
G5	UAT test cases passed / open Critical+High defects	≥ 95% / 0
G6	Open Critical/High security vulnerabilities; PII redaction on seeded test set	0; 100%
G7	Pilot release date and non-labour budget	11 Mar 2026; ≤ 5,500 USD
G8	Story points delivered / Must-have stories delivered	≥ 90% / 100%

ID	Business KPI (after go-live)	Proposed Project Target
K1	AI deflection rate (conversations resolved without human handover)	≥ 35% within 8 weeks of pilot go-live
K2	Response speed	AI answer p95 ≤ 8 s; first human reply to escalations ≤ 2 business hours
K3	Customer feedback on AI answers	≥ 75% thumbs-up within 8 weeks of pilot go-live
K4	Knowledge-base coverage of the 20 most frequent ticket categories	≥ 80% of categories have a published, reviewed article at go-live

11. Governance and authority

Topic	Arrangement
Project Manager authority	Plans and controls the project; owns schedule, risk, issue, dependency and change logs; removes blockers; may re-order work inside a sprint with Product Owner agreement.
Steering	Sponsor, CTO, Support Manager, Product Owner and PM meet at each sprint review (30 minutes, Fridays 14:00).
Decision rights	Product Owner: backlog and priorities. Support Manager: UAT acceptance. Security & Privacy Officer: security/privacy gate. Sponsor: budget, scope of Must-Haves, Go/No-Go.
Change tolerance	PM + Product Owner may defer Should/Could stories with Sponsor informed. Changes to Must-Have scope, budget above +10%, or schedule above +3 working days go to the Sponsor.
Reporting	Weekly status check, one-page status at each sprint review, risk register reviewed weekly.
Escalation	Blocker unresolved for 2 business days: PM escalates to Product Owner; then to Sponsor.

12. Approvals (simulated)

Role	Name	Approval
Executive Sponsor	Marcus Hale (fictional)	Simulated approval - 7 Jan 2026 (M1). No real signature.
Product Owner	Elena Rossi (fictional)	Simulated approval - 7 Jan 2026.
Project Manager	Portfolio author	Prepared the charter.