Project Closure Report

Simulated closure scenario - 13 March 2026

Document	14 - Project Closure Report
Project	AI Customer Support Platform (fictional company: Kestrel Commerce Ltd.)
Prepared by	Project Manager (portfolio author) - self-directed simulation
Status	Simulated closure scenario
Project window	5 Jan - 13 Mar 2026 (simulated project window)

SIMULATED PROJECT - Self-directed portfolio case study. Not professional experience, not client work, not employment. All people, figures and outcomes are fictional.

This is a simulated closure scenario written to show how a closure report is structured. The delivery figures describe an illustrative scenario that continues the simulated status report; they are not real project results. Business KPIs are Proposed Project Targets that could only be measured after a real pilot.

1. Project summary

Project	AI Customer Support Platform - Kestrel Commerce Ltd. (fictional)
Planned window	5 Jan 2026 - 13 Mar 2026
Scenario outcome	Pilot released to production on 11 Mar 2026 (baseline date); closure on 13 Mar 2026.
Scope delivered (scenario)	158 of 161 story points (98%); all 26 Must-Have stories; 3 of 4 Should-Have stories (audit logging deferred).
Methodology	Scrum, five two-week sprints (proposed delivery methodology in a self-directed simulation).

2. Deliverables

Deliverable	Status (scenario)	Notes
Charter, scope, requirements, backlog, UX prototype	Delivered	Baselined at M1-M2.
Knowledge-base admin (create, edit, publish, import)	Delivered	Approximately 262 articles published after audit; 17 flagged articles held back until rewritten.
AI assistant (ingestion, RAG, guardrails, golden set)	Delivered	Three evaluation runs recorded; gate targets reviewed at Go/No-Go.
Customer chat, sign-in, history, feedback	Delivered	US-17 moved from Sprint 3 to Sprint 4 (Decision D1).
Escalation queue, agent console, reply notification	Delivered	Standalone; no helpdesk integration (Won't Have).
Analytics events and manager dashboard	Delivered	Definitions documented, including deflection rate.
Security and privacy controls	Delivered	Hardening, PII redaction, retention and deletion.
Audit logging (US-08, Should)	Deferred	Moved to post-MVP by Product Owner and PM when UAT fixes used Sprint 5 capacity; Sponsor informed. Basic access logs remain.
UAT sign-off, load test, runbook, guides, training	Delivered	UAT closed with 3 Low-severity known issues accepted by the Support Manager.
Production pilot release and hypercare	Delivered	Released 11 Mar; hypercare to 13 Mar.

3. Delivery performance (scenario)

Sprint	Planned SP	Delivered SP	Variance	What moved
Sprint 1	28	25	-3	US-05 completed early in Sprint 2.
Sprint 2	36	34	-2	US-05 in, US-15 out (carried).
Sprint 3	34	36	+2	US-15 carried in; US-17 moved out (D1).
Sprint 4	34	37	+3	US-17 moved in; velocity stabilised after tuning.
Sprint 5	29	26	-3	US-08 deferred.
Total	161	158	-3	98% of planned points; 100% of Must-Have stories.

- Schedule: all nine milestones met on baseline dates in the scenario; no change to the 11 Mar pilot date.
- Budget: approximately 3,800 of 5,500 USD non-labour budget used (simulated figure); contingency unused.
- Scope changes handled: CR-001 Spanish support deferred; US-17 re-sequenced; US-08 deferred. No Must-Have change.
- Risks: R03 (content quality) was the main one realised and was contained by the audit; R01 was controlled through evaluation gates; no privacy or security incident in the scenario.

4. Review of proposed success criteria

ID	Criterion	Proposed Project Target	Closure assessment (scenario)
G1	Golden-set answer accuracy	≥ 85%	Met (assumed in scenario); evidence would be evaluation run #3.
G2	Unsupported answers	≤ 5%	Met (assumed in scenario); same evidence.
G3	Unanswerable questions falling back	≥ 95%	Met (assumed in scenario).
G4	p95 response time at 50 sessions	≤ 8 s, errors < 1%	Met (assumed in scenario); load test report.
G5	UAT pass rate / open Critical+High defects	≥ 95% / 0	Met with 3 Low known issues accepted.
G6	Open Critical/High vulnerabilities; PII redaction	0; 100%	Met (assumed in scenario); security test and redaction test results.
G7	Pilot date and budget	11 Mar 2026; ≤ 5,500 USD	Met.
G8	Points delivered / Must-Have stories	≥ 90% / 100%	Met: 98% / 100%.
K1	Deflection rate	≥ 35% within 8 weeks of go-live	Not measurable at closure; measured weekly by the Support Manager for 8 weeks.
K2	Response speed	AI p95 ≤ 8 s; human reply ≤ 2 business hours	Partly measurable at closure (AI time); human reply measured in pilot.
K3	Positive feedback on AI answers	≥ 75%	Not measurable at closure; dashboard in place.
K4	KB coverage of top 20 categories	≥ 80% at go-live	Met (assumed in scenario) after content audit.

5. Outstanding items

Item	Owner	Due	Notes
US-08 Audit logging (post-MVP backlog)	Product Manager	Next planning cycle	Security-related Should; recommend scheduling early.
US-31 KB version history, US-32 sentiment escalation (Could)	Product Manager	After pilot review	Candidates for the follow-on release.
3 Low-severity known issues from UAT	Backend / Frontend	First post-pilot sprint	Logged and accepted.
Measure K1-K3 for 8 weeks after go-live	Support Manager	8 weeks after 11 Mar	Dashboard is the source; baseline documented.
Weekly review of 20 sampled AI conversations (AIR-09)	Support Manager	Weekly during pilot	Feeds content fixes.
LLM spend monitoring (alerts 70% / 90%)	Support Operations	Ongoing	Handover from DevOps.
KB content ownership and monthly review	Knowledge & Content Lead	Monthly	Mitigates R17.

6. What went well, what could improve, recommendations

What went well (scenario)	What could be improved	Recommendations
Grounded answers with a safe fallback built confidence with stakeholders. Change control kept scope stable. Evaluation in the Definition of Done caught quality regressions early. One-page status and decision table sped up decisions.	Provider terms and SME time should have been booked as dependencies in Week 1. A Should-Have security story (audit logging) was planned too late to survive. Load test left only three days to react.	Start vendor/legal review and SME booking at kickoff. Schedule security-relevant Should items before the final sprint. Run a smoke test one sprint before the load test. Use the first two sprints' velocity to re-forecast, not the plan.

Detailed lessons are in 15_Lessons_Learned.pdf.

7. Handover and acceptance (simulated)

- Operations: Support Operations owns the platform; DevOps provides the runbook, dashboards and alerts.
- Content: Knowledge & Content Lead owns the knowledge base and the monthly review.
- Backlog: post-MVP backlog (US-08, US-31 to US-35) handed to the Product Owner.
- Acceptance: simulated Sponsor acceptance on 12 Mar 2026; no real signature.

8. Limitations of this report

This closure report is part of a self-directed simulation. No real users, data or AI measurements exist behind it, so the assessments marked 'assumed in scenario' show how a PM would structure the evidence, not an achieved result.